

# 機械学習特論

～ 理論とアルゴリズム ～

(Clustering: k-means and EM)

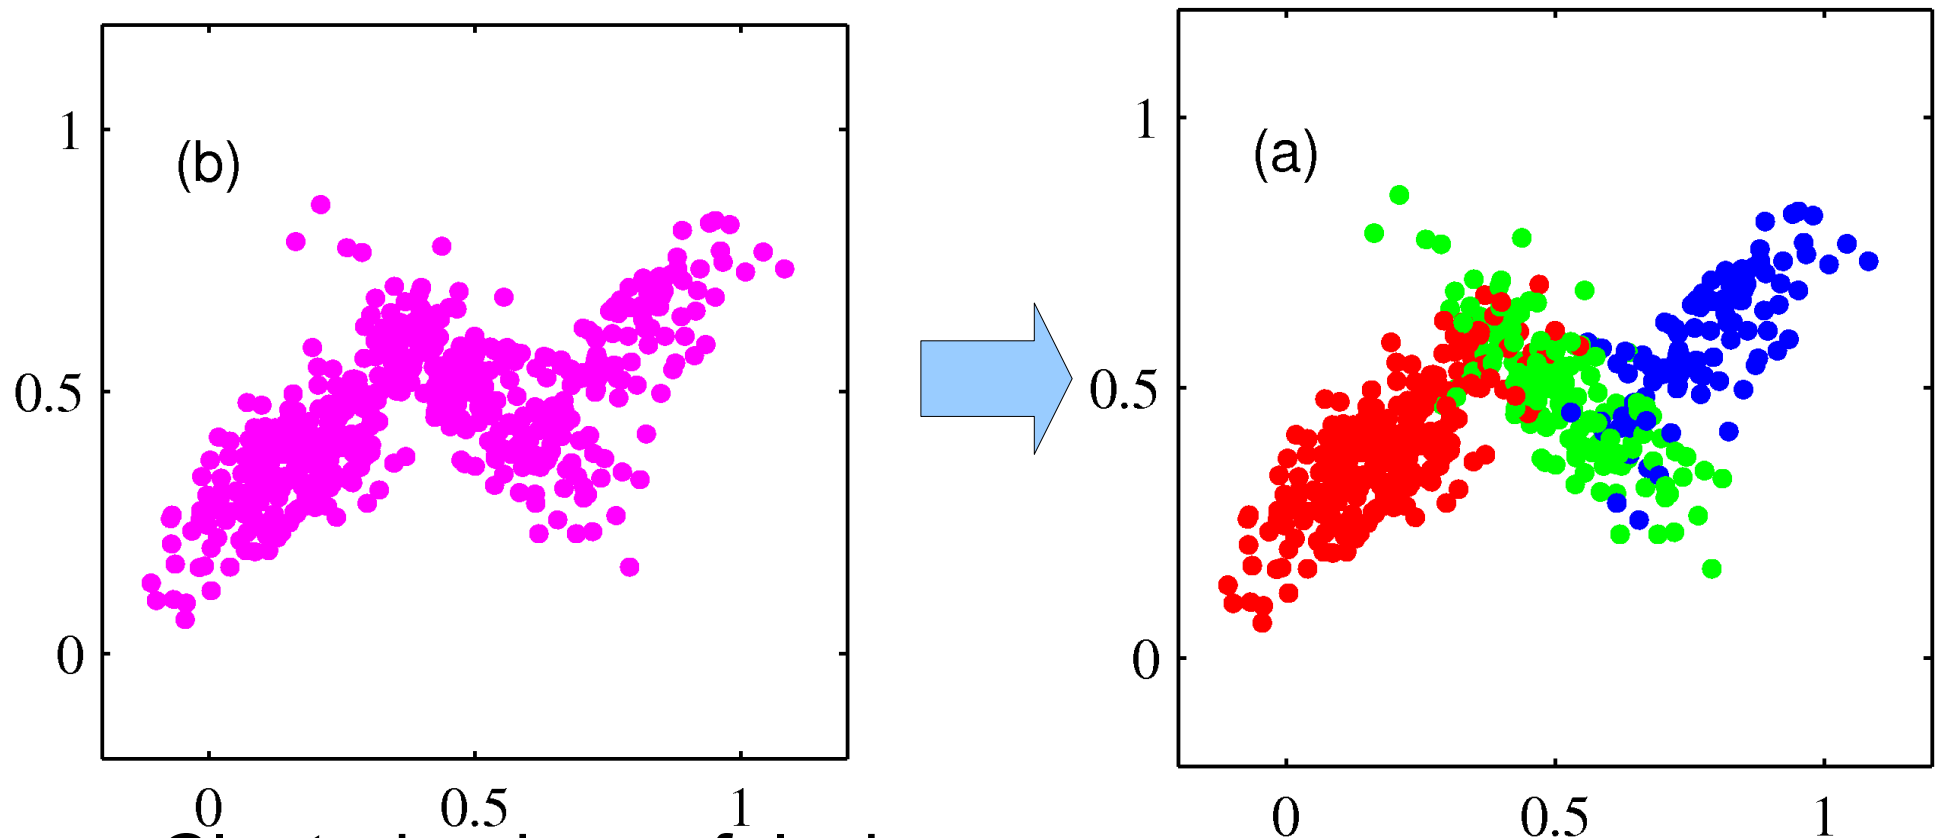
講師：西郷浩人

# Supervised and Unsupervised Learning

- Supervised Learning( 教師あり学習 )
  - Classification( 分類 )
    - Given observed data(explanatory variables)  $X \in \mathbb{R}^{n \times p}$  predict a category  $y \in \mathbb{N}^n$
  - Regression( 回帰 )
    - Given observed data  $X \in \mathbb{R}^{n \times p}$  predict real-valued response  $y \in \mathbb{R}^n$
- Unsupervised Learning( 教師なし学習 )
  - Clustering
    - Given observed data  $X \in \mathbb{R}^{n \times p}$  generate categories and classify datum into one of them

# Clustering( クラスタリング )


- The goal is to classify data points into categories. In this case, into 3 categories.



- Clustering is useful when we have limited knowledge about the data. (e.g., we don't even know the names of fish.)

# Mixture of Gaussians

- Superimposing several gaussian(normal) distribution enables us to represent various types of distributions.
- Below we assume that our data consist of mixture of gaussians (especially in the EM case)




# Clustering algorithms

- k-means
  - Simple
  - Often gives us satisfactory performance
  - Only estimates means of clusters
- Expectation Maximization
  - More complicated
  - Estimates means and covariances of clusters


# k-means (k- 平均法 )

- Algorithm k-means
  - Input: data points ( $n$ ), number of clusters ( $k$ )
  - Outputs: assignment of data points to each cluster
  - Initialize:
    - set cluster centers randomly
  - Iteration:
    - assign each data point to the nearest cluster centers
    - renew cluster centers by taking mean of data points.

N:50 k:2 Iteration: 1




N:50 k:2 Iteration: 2






N:50 k:2 Iteration: 3



N:50 k:2 Iteration: 4



# Ex. 1

- Download today's data, and observe the behavior of k-means algorithm.
- Inputs
  - X: data (explanatory variables)
  - k: number of clusters
- Observe the behavior of the algorithm by changing k.
  - Usage: `[estimated_mu] = kmeans(X,k)`

# About the data

- $X$  ( 50 x 2 matrix )
  - The former 25 data points are generated from gaussian distribution with mean (1,1)
  - The latter 25 data points are generated from gaussian distribution with mean (-1,-1)
- Generated by gen2Ddata.m

## gen2Ddata.m

```
function [X] = gen2Ddata(N,meanX,meanY)
% generate 2D data from gaussian distribution
% Inputs
% N: number of data points
% meanX: 2-dimensional vector specifying two x coordinates
% meanY: 2-dimensional vector specifying two y coordinates
%
% example: meanX=[-1 1], meanY=[-1 1], X=gen2Ddata(50,meanX,meanY);

var = 1;
mid = floor(N/2);
for i=1:mid
    x(i)=normrnd(meanX(1),var);
    y(i)=normrnd(meanY(1),var);
end
for i=mid+1:N
    x(i)=normrnd(meanX(2),var);
    y(i)=normrnd(meanY(2),var);
end
X=[x' y'];
```

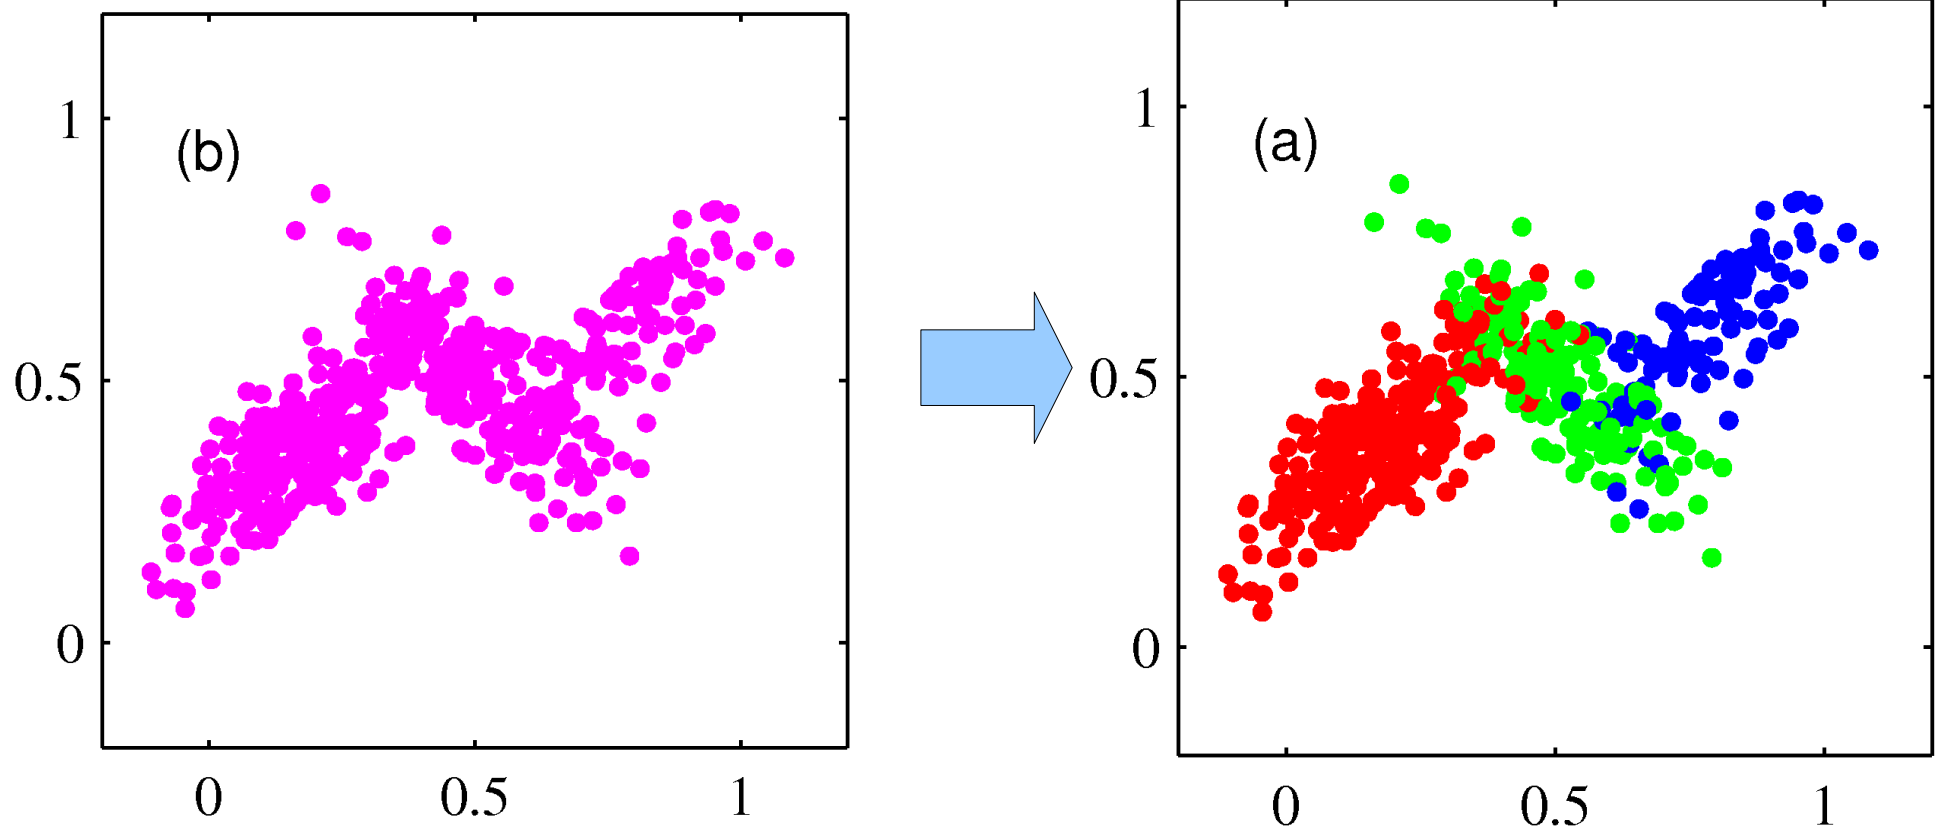
Note: Your environment may not have function 'normrnd'. If so, you can search for 'normrnd.m' on the web, and download it to your current directory

# Ex1 continued.

- Algorithm chooses cluster centers randomly, therefore the results you obtained can be different each time even though the input does not change.
- What do you observe by increasing the number of clusters to 2, 3 ?

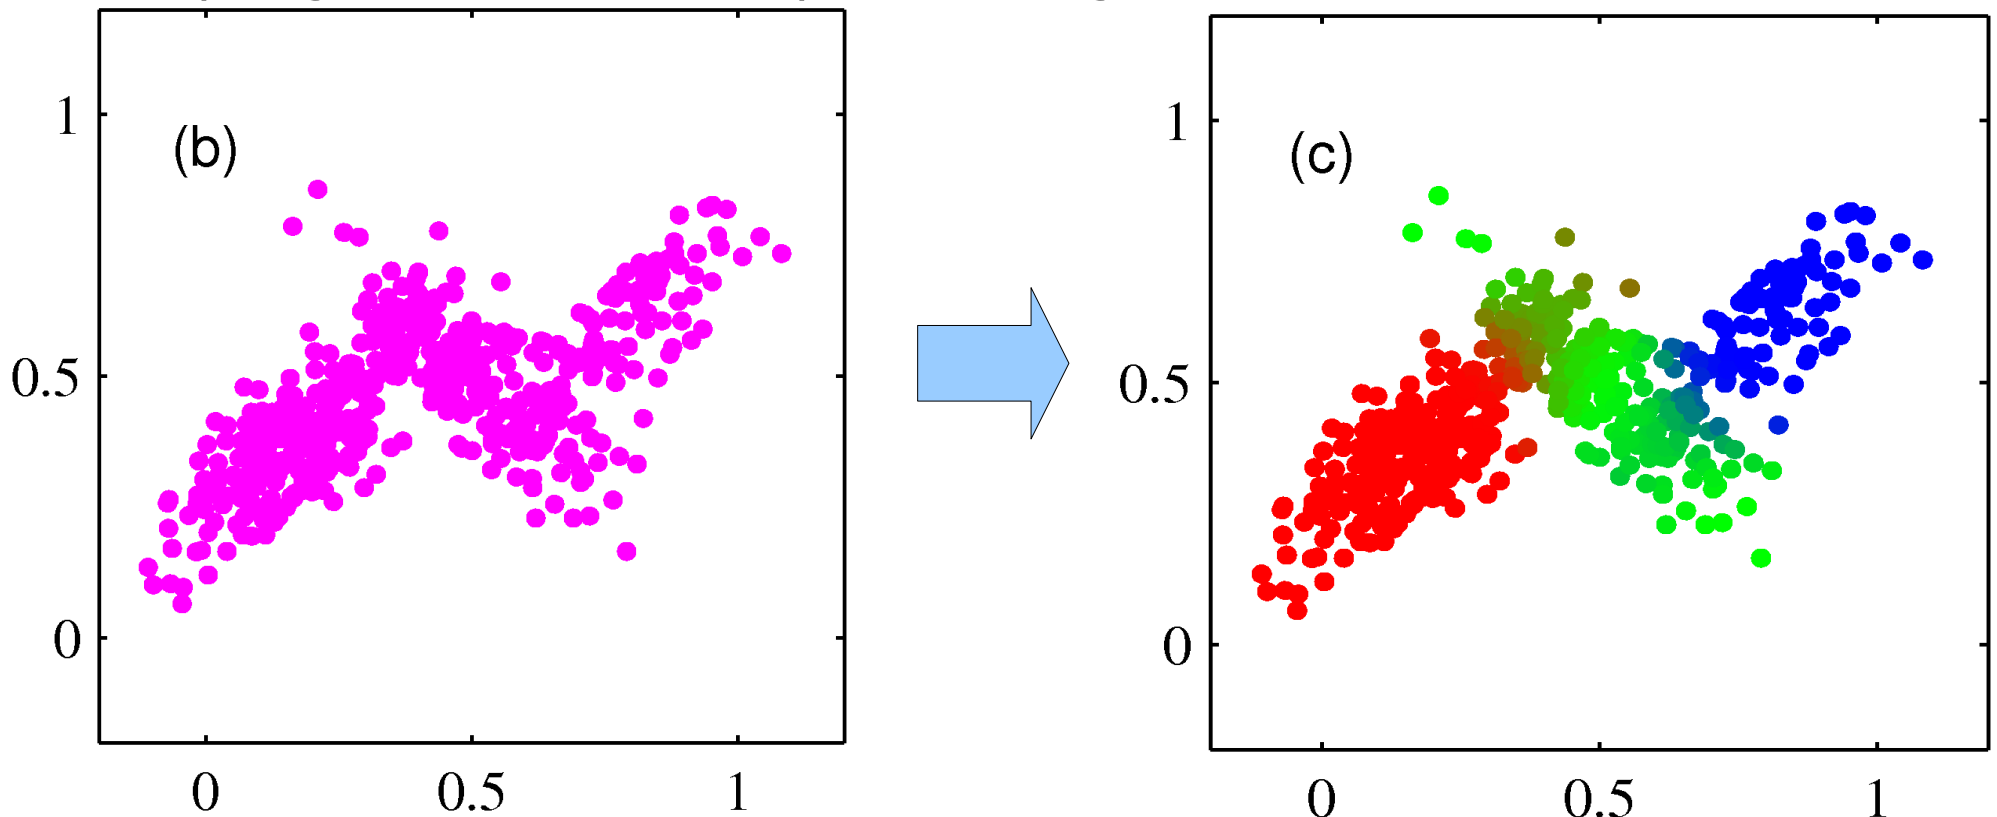
# K-means as *hard* clustering

- Each data point is either assigned to cluster(s) or not.
- In this sense, k-means is a hard clustering algorithm.



# Soft clustering

- Now we allow each data point to be shared by multiple clusters with arbitrary mixing ratio, then membership probability can take either  $[1, 0, 0]$ ,  $[0.5, 0.5, 0]$ ,  $[0.3, 0.3, 0.3]$  etc.
- Notice the data points in between red and green (or green and blue) in the figure below.





# Preparation

- Let us define mixture weights (prior probability) as proportion of each clusters that sum to one:

$$\sum_{k=1}^K \alpha_k = 1$$

- Then probability of each data point is weighted sum of probabilities

$$p(\mathbf{x}_i) = \sum_{k=1}^K \alpha_k p_k(\mathbf{x}_i | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

where

$$p_k(\mathbf{x}_i | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\boldsymbol{\Sigma}_k|^{1/2}} \exp\left(\frac{-1}{2} (\mathbf{x}_i - \boldsymbol{\mu}_k)' \boldsymbol{\Sigma}_k^{-1} (\mathbf{x}_i - \boldsymbol{\mu}_k)\right)$$

# Posterior probability

- Then posterior probability of a point belonging to class  $k$  is determined by Bayes' rule.
  - Mixture weights  $\alpha$  can be thought as priors.

$$w_{i,k} = p(C=k|\mathbf{x}_i, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) = \frac{\alpha_k p_k(\mathbf{x}_i|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{k=1}^K \alpha_k p_k(\mathbf{x}_i|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}$$

## Bayes' rule

$$p(y|x) = \frac{p(y)p(x|y)}{p(x)} = \frac{p(y)p(x|y)}{\sum_y p(y)p(x|y)}$$

# Maximumizing Loglikelihood

$$\begin{aligned} l &= \log \prod_{i=1}^N p(\mathbf{x}_i) \\ &= \log \prod_{i=1}^N \sum_{k=1}^K \alpha_k p_k(\mathbf{x}_i | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \\ &= \sum_{i=1}^N \log \sum_{k=1}^K \alpha_k p_k(\mathbf{x}_i | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \end{aligned}$$

- Parameters that maximize the above log-likelihood is obtained by solving the following set of equations.

$$\frac{\partial l}{\partial \boldsymbol{\alpha}_k} = 0 \quad \longleftrightarrow \quad \alpha_k^{new} = \frac{1}{N} \sum_{i=1}^N w_{i,k}$$

$$\frac{\partial l}{\partial \boldsymbol{\mu}_k} = 0 \quad \longleftrightarrow \quad \boldsymbol{\mu}_k^{new} = \left( \frac{1}{\sum_{i=1}^N w_{i,k}} \right) \sum_{i=1}^N w_{i,k} \mathbf{x}_i$$

$$\frac{\partial l}{\partial \boldsymbol{\Sigma}_k} = 0 \quad \longleftrightarrow \quad \boldsymbol{\Sigma}_k^{new} = \left( \frac{1}{\sum_{i=1}^N w_{i,k}} \right) \sum_{i=1}^N w_{i,k} (\mathbf{x}_i - \boldsymbol{\mu}_k^{new})(\mathbf{x}_i - \boldsymbol{\mu}_k^{new})'$$

# Wait a moment..

$$\alpha_k^{new} = \frac{1}{N} \sum_{i=1}^N w_{i,k}$$

$$\mu_k^{new} = \left( \frac{1}{\sum_{i=1}^N w_{i,k}} \right) \sum_{i=1}^N w_{i,k} \mathbf{x}_i$$

$$\Sigma_k^{new} = \left( \frac{1}{\sum_{i=1}^N w_{i,k}} \right) \sum_{i=1}^N w_{i,k} (\mathbf{x}_i - \mu_k^{new})(\mathbf{x}_i - \mu_k^{new})'$$

- Solving the above equations is not trivial, since they all depend on unknown posterior probability  $w_{i,k}$ , which recursively depend on  $\alpha_k, \mu_k, \Sigma_k$

$$w_{i,k} = \frac{\alpha_k p_k(\mathbf{x}_i | \mu_k, \Sigma_k)}{\sum_{k=1}^K \alpha_k p_k(\mathbf{x}_i | \mu_k, \Sigma_k)}$$

# EM (Expectation Maximization) algorithm for gaussian mixtures

$$w_{i,k} = \frac{\alpha_k p_k(\mathbf{x}_i | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{k=1}^K \alpha_k p_k(\mathbf{x}_i | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}$$

$$\alpha_k^{new} = \frac{1}{N} \sum_{i=1}^N w_{i,k}$$

$$\boldsymbol{\mu}_k^{new} = \left( \frac{1}{\sum_{i=1}^N w_{i,k}} \right) \sum_{i=1}^N w_{i,k} \mathbf{x}_i \quad \longleftarrow \quad \text{Weighted mean vector (p x 1)}$$

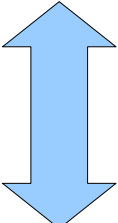
$$\boldsymbol{\Sigma}_k^{new} = \left( \frac{1}{\sum_{i=1}^N w_{i,k}} \right) \sum_{i=1}^N w_{i,k} (\mathbf{x}_i - \boldsymbol{\mu}_k^{new})(\mathbf{x}_i - \boldsymbol{\mu}_k^{new})' \quad \longleftarrow \quad \text{Weighted covariance matrix (p x p)}$$

- So we initialize  $\alpha_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k$  by some arbitrary values, and compute  $w_{i,k}$
- Then recompute  $\alpha_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k$ , and iterates...

# EM (Expectation Maximization) algorithm for gaussian mixtures

E-step

$$w_{i,k} = \frac{\alpha_k p_k(\mathbf{x}_i | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{k=1}^K \alpha_k p_k(\mathbf{x}_i | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}$$

  
M-step

$$\alpha_k^{new} = \frac{1}{N} \sum_{i=1}^N w_{i,k}$$

$$\boldsymbol{\mu}_k^{new} = \left( \frac{1}{\sum_{i=1}^N w_{i,k}} \right) \sum_{i=1}^N w_{i,k} \mathbf{x}_i$$

$$\boldsymbol{\Sigma}_k^{new} = \left( \frac{1}{\sum_{i=1}^N w_{i,k}} \right) \sum_{i=1}^N w_{i,k} (\mathbf{x}_i - \boldsymbol{\mu}_k^{new})(\mathbf{x}_i - \boldsymbol{\mu}_k^{new})'$$

- In the framework of EM algorithm, computation of  $\alpha_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k$  is called M-step, and that of  $w_{i,k}$  is called E-step.

# Convergence of EM algorithm

- Convergence of the algorithm is checked by observing the change of log-likelihood.
- **EM algorithm for gaussian mixtures only converges to its local maximum, but not to global maximum.**
- For searching global maximum..
  - Algorithm can be run again and again using different initial parameters.

## Ex. 2





- Download "EM.m" and observe the behavior of the EM algorithm. Read the code by yourself.
- Inputs:
  - X: data
  - k: number of clusters
  - usage:  $[\alpha, \mu, \sigma, w] = \text{EM}(X, k)$



## Ex. 2 continued.

- How do you interpret the outputs of EM.m: alpha, mu, sigma, w. Do they look relevant ?
  - Hint: `imagesc(w)`;
    - w contains posterior probabilities for class labels.
- Generate your own data and try EM algorithm on it !
  - Use of `gen2Ddata.m`
    - `n=50, mu1=[1 1], mu2=[-1 -1]`
    - `[X]=gen2Ddata(n,mu1,mu2)`

# Results are different each time



## Ex. 3

- Compare K-means and EM in terms of the following points.
  - Speed, accuracy, mean, variance
- For measuring the speed, you should turn visualization off by setting the third input to zero, and use 'tic' and 'toc'
  - `tic, [mu]=kmeans(X,k,0), toc`
  - `tic, [alpha mu sigma w]=EM(X,k,0), toc`

# How to set number of clusters ?

- In clustering, *how to set the number of clusters* is an open problem.
- However, EM algorithm returns maximized likelihood, thus one can make use of information criteria such as
  - $AIC = -2 \log \text{likelihood} + 2p$ 
    - where  $p$  is the number of parameters
  - $BIC = -2 \log \text{likelihood} + p \log n$ 
    - where  $n$  is the number of data points.

# Appendix

# Tips: Number of clusters

- On today's data, EM algorithm was run 100 times for each k, and computed average AIC and BIC. Values inside parentheses show standard deviation.

k	1	2	3	4	5
AIC	435.7(0.00)	434.4 (3.29)	435.5 (5.02)	437.4 (9.20)	443.4 (11.7)
BIC	443.4 (0.00)	451.7 (3.29)	464.2 (5.02)	479.5 (9.20)	500.8 (11.7)

- In this case AIC did a better job, but BIC underestimated the number of clusters.
- A script used for this experiment is included, and named as "checkNumCluster.m"